Context Aware Music Recommendation System

Pavan Kalyan Reddy Cherupally
Arizona State University
Tempe, AZ, USA
pcherup1@asu.edu

Tanmai Mukku
Arizona State University
Tempe, AZ, USA
tmukku@asu.edu

Spoorthi Uday Karakaraddi
Arizona State University
Tempe, AZ, USA
sudaykar@asu.edu

Sai Siddharth Vemuri
Arizona State University
Tempe, AZ, USA
svemur13@asu.edu

ABSTRACT

Music recommendation systems play a pivotal role in the digital music industry, significantly enhancing user experience by personalizing music selections. However, the majority of existing systems predominantly focus on leveraging user preferences and song metadata, neglecting the critical influence of contextual information on music choice. Our project, "Context-Aware Music Recommendation System," seeks to bridge this gap by incorporating a variety of contextual factors, such as the user's emotional state, physical environment, and temporal variables, into the recommendation process. This report presents the initial progress of our project, including a comprehensive review of related work, the development and preprocessing of our datasets, and a preliminary methodology for context incorporation. We discuss the challenges encountered in integrating contextual data and outline our proposed solutions and future work directions. Our findings suggest that contextual information significantly enhances recommendation relevance, setting a foundation for further exploration in the pursuit of a more nuanced and user-centered approach to music recommendation.

KEYWORDS

Recommendation Systems, Collaborative Filtering, Factorization Machines

1 INTRODUCTION

Music recommendation systems are crucial in the digital age, where the abundance of available content can overwhelm users. These systems aim to personalize the user experience by suggesting songs, artists, or playlists that align with individual tastes and preferences. The evolution of recommendation systems has led to the development of various approaches, each with its strengths and limitations.

1.1 Collaborative Filtering

Collaborative Filtering (CF) is one of the most widely used techniques in recommendation systems. It makes predictions about a user's interests by collecting preferences from many users. The underlying assumption is that if users agreed in the past, they would agree in the future. Despite its popularity, CF suffers from cold start and sparsity issues, limiting its effectiveness when new users or items have insufficient interaction data.

1.2 Content-Based Filtering

Content-Based Filtering approaches recommend items by comparing the content of the items and a user profile. In the context of music, this could involve analyzing the audio features of tracks, such as tempo, genre, and harmony, or lyrical content to predict user preferences. While addressing some of the CF limitations, content-based methods can lead to a lack of diversity in recommendations, as they rely solely on item attributes.

1.3 Hybrid Methods

Hybrid methods combine collaborative and content-based filtering to leverage the strengths of both. By integrating diverse data sources, hybrid systems can provide more accurate and personalized recommendations, mitigating the limitations of individual approaches. These methods have shown significant promise in improving recommendation quality and user satisfaction.

1.4 Context Aware and 3D Algorithms

The advent of **deep learning** and **context-aware algorithms** has introduced a new dimension to recommendation systems. These advanced techniques consider not only the historical preferences of users and the characteristics of music tracks but also the contextual information surrounding each listening session. In our exploration, we are particularly interested in the potential of **Factorization Machines** for their ability to model interactions among users, items, and contextual factors efficiently.

This paper is organized as follows: Section 2 reviews related work in the field of music recommendation systems, highlighting the evolution from traditional methods to advanced, context-aware approaches. Section 3 presents an exploratory data analysis (EDA) of the datasets employed in our study. Section 4 details the methodology, with a focus on our considerations for employing Factorization Machines. Section 5 describes the experiments conducted, followed by Section 6, which presents the results. Section 7 discusses the implications of our findings, and Section 8 concludes the paper with reflections on our research and future directions.

2 RELATED WORK

Recent advancements in music recommendation systems have significantly emphasized the integration of contextual information and the application of innovative machine learning techniques to enhance personalization and accuracy. A growing body of research explores a variety of contextual data sources, such as playlist names, user behaviors, and audio-lyric features, to better understand user preferences in specific situations. Studies by Pichl et al. [6] and Baxter et al. [2] highlight the importance of contextual clues from playlist names and song attributes in achieving more accurate user preference predictions than traditional collaborative filtering methods. Similarly, Zamani et al. [10] and another study by Pichl et al. [5] emphasize the significance of situational context integration, indicating that multi-context-aware systems can considerably enhance recommendation relevance and accuracy.

Furthermore, the effectiveness of machine learning algorithms in extracting deep features from music content and user interaction data is widely recognized. Chang et al. [3] demonstrate the potential of combining audio and lyric features with reinforcement learning for dynamic recommendations that adapt over time. Moreover, Yong Zheng and David Wang [11] propose a multi-criteria ranking framework that utilizes user criteria, such as click-through rates and ratings, showcasing that sophisticated data analysis techniques can yield high-quality recommendations.

Nonetheless, the scalability and generalizability of these approaches remain subjects of debate. Adomavicius and Tuzhilin [1] provide a comprehensive framework for context-aware recommendations, acknowledging the necessity for broader experimentation and more concrete implementations across different domains. The dependency on specific types of contextual information also raises concerns regarding the depth of personalization and the systems' ability to capture accurate situational contexts. Elbir and Aydin [4] discuss how the success of a recommendation system based on Convolutional Neural Networks hinges on the noise within the training data, underscoring the need for models that account for a broader range of contextual factors.

In addition to exploring these approaches, our research leans towards the application of Factorization Machines (FMs), introduced by Rendle [8], for their capability to model interactions among users, items, and contextual factors efficiently, even with sparse data. Adomavicius et al. [1] suggest that 3D algorithms such as Context-Aware Recommender Systems (CARS) can be simplified to a 2D recommendation algorithm with a contextual filtering step, illustrating the potential for polynomial regression models and FMs to achieve contextual modeling in music recommendations.

This body of work forms the foundation for our exploration of context-aware music recommendation systems, acknowledging the ongoing debates and suggesting a direction towards innovative solutions that balance technological sophistication with practical applicability.

3 EXPLORATORY DATA ANALYSIS

3.1 Dataset Overview

We use the #nowplaying-RS [7] dataset for this project. It consists of both context and content based features along with timestamps of listening events. This dataset is structured across three distinct tables, each providing valuable insights into user behavior and preferences. These datasets comprise various aspects of music streaming sessions from Spotify, including sentiment scores derived from Twitter hashtags, track features, and user interaction data. The three data tables are as follows:

- **User Track Hashtags With Timestamps**: It contains basic information about each listening event such as user_id, track_id, hashtag and created_at.
- **Content Features**: It encompasses a comprehensive set of features, incorporating both song characteristics (such as valence, loudness, tempo, etc.) and user-related information, including location, language.
- **Sentiment values**: It contains the sentiment features of all hashtags collected from four different sentiment dictionaries: AFINN, Opinion Lexicon, Sentistrength Lexicon and vader.

3.2 Data Preprocessing

3.2.1 Data Loading and Cleaning. The datasets were loaded and inspected for dimensions and missing values. Preprocessing steps included:

- Converting data types for numeric columns that were mistakenly encoded as objects.
- Detecting and handling missing values by either filling with averages or removing rows with null values.
- Dropping unnecessary columns that did not contribute valuable information for the analysis or modeling processes.

3.2.2 Sentiment Analysis Dataset. The sentiment analysis dataset contained various sentiment scores associated with specific hashtags. The following preprocessing steps were applied:

- Renaming columns for clarity.
- Dropping columns with redundant or no valuable information.
- Filling missing values with column averages where applicable.

3.2.3 Dataset Size Reduction. The dataset comprises over 9 million Listening Events (LEs), which was reduced to 2.5 million to accommodate computational resource constraints. This reduction involved:

- Excluding LEs for songs listened to less than 50 times.
- Filtering out data not associated with UK or US time zones.
- Removing entries without a vader sentiment score.

3.3 Feature Overview

The analysis utilizes features such as sentiment_score, instrumentalness, danceability, energy, loudness, tempo, acousticness, valence, mode, and key to understand music preferences and listening behaviors.

3.4 Exploratory Data Visualization

Visualizations were created to explore the distribution of sentiment scores and track features. Key insights include:

- A wide range of sentiment scores, indicating diverse emotional contexts associated with music listening experiences.

- The distribution of track features such as danceability, energy, and acousticness, highlighting the variability in musical content.

3.4.1 Correlation Matrix. Pearson's correlation coefficient was used to estimate correlations among musical features, revealing:

- A positive correlation between energy and loudness, suggesting that tracks with high energy typically exhibit higher loudness levels.
- A negative correlation between energy and acousticness, indicating that songs with higher energy are generally less acoustic.

Notably, the 'mode' and 'key' features, being categorical, do not imply any correlation through this analysis.

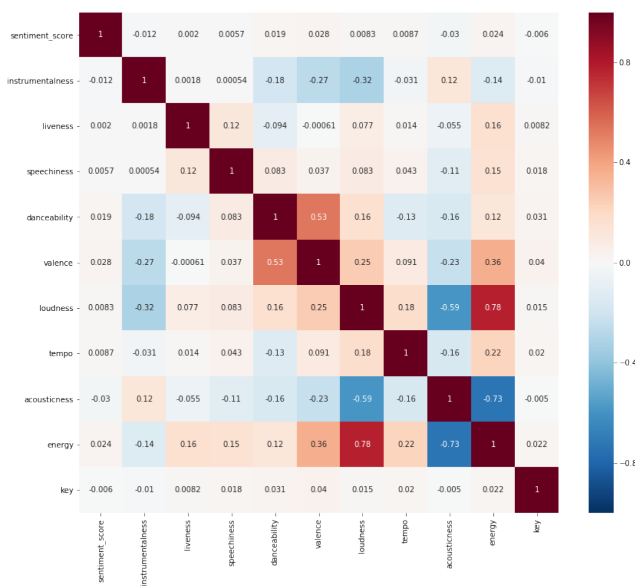

Figure 1: Estimating correlations among the musical features by visualizing correlation matrix after computing Pearson's correlation coefficient

3.4.2 Context-Based Recommendations Insight. Visualization of the danceability feature against mood (sadness and happiness) demonstrates:

- Sad songs (low valence) tend to have lower danceability values, whereas happy songs (high valence) show the opposite trend. This aligns with intuitive expectations regarding the emotional content of music and its impact on danceability.
- Extremes in acousticness are more often associated with sad songs, while the loudness feature inversely correlates with sadness, suggesting quieter songs are more likely to evoke sadness.
- Tempo analysis further supports that slower-paced songs are more often perceived as sad compared to their faster counterparts.

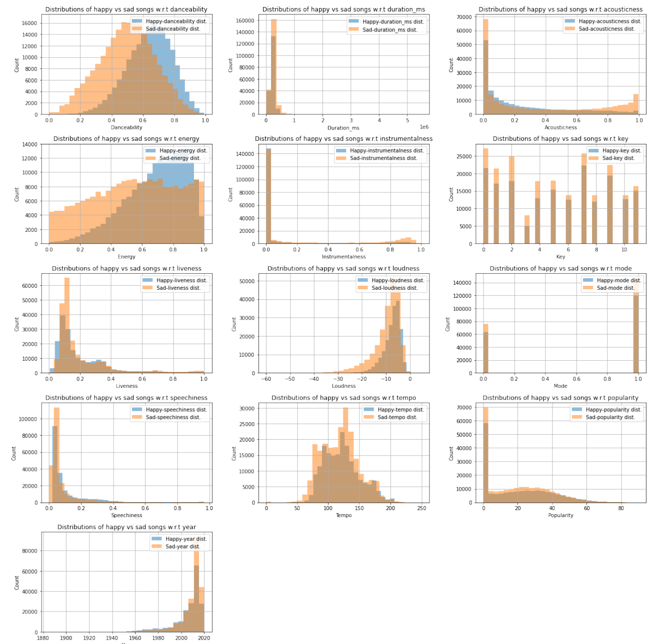

Figure 2: Plots throwing light on context (in this case, mood - happy or sad) based recommendation

3.5 Initial Data Insights

From the exploratory data analysis, several initial insights were gathered:

- The significant impact of sentiment scores on music recommendation, emphasizing the need for context-aware recommendation systems.
- The potential for deep learning models to capture complex patterns in music preference and listening behavior.

The exploratory data analysis provided foundational insights into the datasets, guiding the subsequent modeling efforts. By understanding the data's characteristics, the team can better tailor the recommendation system to incorporate user sentiment and track features effectively.

4 METHODOLOGY

The methodology of our Context-Aware Music Recommendation System centers around the use of Factorization Machines (FMs) [8], which are advanced machine learning models capable of capturing complex interactions between a large number of features in a highly efficient manner. This section outlines the rationale behind choosing FMs, the specific configurations employed, and the integration of diverse data features crucial for enhancing the personalization of music recommendations.

4.1 Factorization Machines

Factorization Machines (FMs) are a class of machine learning algorithms that are particularly adept at making predictions in situations with sparse data. They operate by factorizing the feature vectors, allowing them to capture all interactions between variables despite data sparsity.

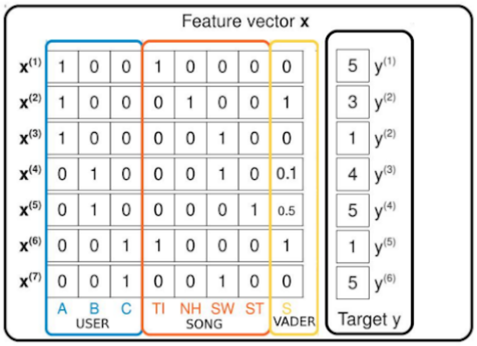

Figure 3: Illustration of a feature vector x utilized by Factorization Machines, with corresponding targets y. The vectors include user identifiers (A, B, C), song identifiers (TI, NH, SW, ST), and sentiment scores (VADER).

As shown in Figure 3, a feature vector x in an FM is typically composed of binary representations for categorical variables, such as user and item IDs, as well as real-valued features such as sentiment scores. In the context of music recommendation systems, these feature vectors might include indicators for the active user (e.g., User A), the item (e.g., Song TI), and other features such as the sentiment derived from VADER scores.

FMs are particularly powerful because they can model the interaction between any two features in the vector. For instance, if User A has a positive interaction with Song TI, and Song TI has a high sentiment score, FMs can effectively capture and utilize this interaction to predict User A's preference for similar songs with high sentiment scores.

This ability to handle a wide range of feature types and capture interactions makes FMs an excellent choice for recommendation systems where the data is often highly dimensional and sparse. By leveraging the factorization concept, FMs can perform efficient computations even in the presence of a vast number of features, which is a common challenge in real-world recommendation systems.

4.2 Choice of Factorization Machines

Factorization Machines are chosen for their unique ability to model all interactions between features within sparse datasets, which is a common scenario in music recommendation systems due to the vast amount of data and the sparse nature of user interactions. Unlike standard machine learning models that struggle with feature sparsity and high dimensionality, FMs effectively handle these challenges by using a factorized parameterization approach. This enables them to estimate reliable parameters even when data is missing extensively, which is a critical advantage in our context.

4.3 Configuration of Factorization Machines

Our implementation of FMs, based on the libFM tool [9], is configured to capture two-way interactions among features. The dimensionality of these interactions is set to five to balance complexity and performance, and we conduct ten iterations of learning to optimize the model's accuracy and robustness. This setup is designed to refine the system's ability to predict user preferences with high precision, leveraging the rich contextual and content-based features available in the #nowplaying-RS dataset.

4.4 Integration of Contextual and Content Features

The methodology also involves a detailed integration strategy for contextual and content features, which are crucial for enhancing the accuracy and relevance of the recommendations. These features include:

- **User and Track Metadata:** Basic identifiers such as User ID and Track ID are included to establish a foundational model of interactions.
- **Enhanced Content Features:** Attributes like valence and tempo of tracks, which represent the emotional and rhythmic characteristics of music, are integrated to capture the subtle nuances that influence user preferences.
- **Contextual Features:** Information such as the time of the tweet and the user's timezone are used to contextualize the recommendations based on when and where users listen to music.
- **Sentiment Analysis:** Sentiment scores derived from hashtags in user tweets provide insights into the emotional states influencing user preferences, adding another layer of personalization to the recommendations.

4.5 Experimental Setup

For the evaluation of our system, the dataset is split into training and test sets based on timestamps, ensuring that the model is tested on future data, simulating a real-world scenario. Negative samples are generated to complement the inherently positive dataset, enhancing the robustness and discriminative power of the model during training. This setup not only tests the effectiveness of our Factorization Machine model in a controlled environment but also mirrors potential real-world applications where the prediction of future user preferences is paramount.

The methodology adopted in this project not only leverages the advanced capabilities of Factorization Machines but also integrates a variety of features that are essential for understanding and predicting the complex preferences of users in real-time and context-aware settings.

5 EXPERIMENTS

The experimental framework is designed to rigorously evaluate the performance of our Context-Aware Music Recommendation System, using the Factorization Machines model equipped with an array of contextual and content features. This section details the setup, methodologies, and evaluation metrics used to assess the system's effectiveness in delivering personalized music recommendations.

5.1 Experimental Setup

Our experiments were conducted using the #nowplaying-RS dataset, which includes a rich set of user interactions, track features, and contextual information. The dataset was divided into training and testing sets based on timestamps to ensure that the model predicts future user preferences, reflecting a realistic usage scenario. Specifically, data from January 1 to September 30 served as the training set, while data from November 1 to December 23 was used for testing. The intervening period in October was optionally used to create a validation set to fine-tune model parameters.

5.2 Generation of Negative Samples

Since the #nowplaying-RS dataset inherently contains only positive examples (user-track interactions that occurred), it is crucial to generate negative samples to provide a balanced view of user preferences. For each positive example in the training and test sets, we introduced nine negative samples. These samples were generated using two methods:

- **Random Population (POP RND):** Negative samples were selected randomly from tracks not previously listened to by the user, ensuring that these tracks do not reflect the user's historical preferences.
- **User-Based Population (POP USER):** Negative samples consisted of tracks that the user had listened to previously, but in different contexts, presenting a more challenging scenario where the model must discern subtle context shifts to make accurate recommendations.

5.3 Model Training and Iterations

The Factorization Machine model was trained using libFM, with a focus on capturing two-way interactions among the features. The model's dimensionality was set to five, optimizing computational efficiency without sacrificing the ability to model complex interactions. The training involved ten iterative passes over the dataset, allowing the model to refine its predictions based on the increasingly complex interplay of user preferences and contextual information.

5.4 Evaluation Metrics

The primary metric used to evaluate the effectiveness of the recommendation system was the Mean Reciprocal Rank (MRR). This metric is particularly suited for scenarios where the goal is to rank a set of recommended items such that the item of interest (i.e., the track the user is most likely to enjoy) ranks as high as possible:

$$MRR = \frac{1}{|Q|} \sum_{i=1}^{|Q|} \frac{1}{\text{rank}_i} \tag{1}$$

where $|Q|$ is the number of queries, and rank_i is the rank position of the correct item for the i-th query. The MRR provides a clear measure of the model's ability to not only identify relevant items but also to prioritize them effectively at the top of the recommendation list.

The experiments conducted were designed to comprehensively evaluate the capability of our Context-Aware Music Recommendation System under various scenarios and using robust metrics. The combination of rigorous experimental design, careful negative sample generation, and precise metrics ensures that the findings are both valid and actionable, providing clear pathways for enhancing the recommendation engine further.

6 RESULTS

The performance of the Context-Aware Music Recommendation System was quantitatively evaluated through the Mean Reciprocal Rank (MRR) for each method, focusing on the context-aware settings of POP RND and POP USER. The MRR values provide insight into the system's ability to prioritize the most relevant track within a list of recommendations. The following table summarizes the MRR scores obtained in the evaluation:

Method	POP RND MRR	POP USER MRR
1 [Base]	0.7755	0.3906
2 [Base+Valence]	0.8010	0.6325
3 [Base+Tempo]	0.7241	0.6581
4 [Base+Hashtag]	0.7515	0.7814
5 [Base+Sentiment]	0.7137	0.8854

Table 1: MRR scores for context-aware recommendation methods in POP RND and POP USER settings.

Figure 4 illustrates the MRR scores for the respective methods and settings, providing a visual comparison of the system's performance across different feature integrations.

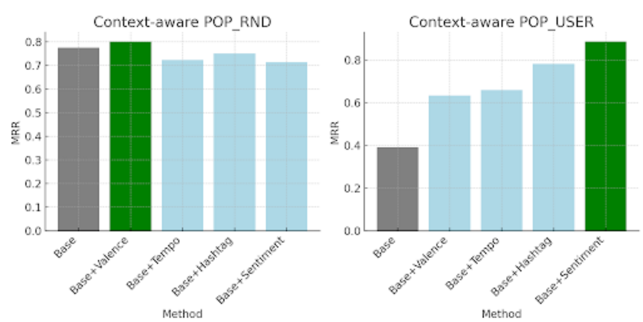

Figure 4: Mean Reciprocal Rank (MRR) for context-aware recommendation methods in POP RND and POP USER settings.

These scores reflect the system's capacity to correctly rank the user's preferred track within an assorted list, with each method introducing a distinct set of features to the Factorization Machine model.

7 DISCUSSION

The evaluation of the Context-Aware Music Recommendation System using Factorization Machines (FMs) and the #nowplaying-RS dataset revealed significant findings concerning the effectiveness of various feature integrations. The Mean Reciprocal Rank (MRR) served as the primary metric for assessing the model's ability to rank a user's preferred track highest among a set of ten recommendations.

7.1 Feature Impact on MRR

The baseline method, which utilized only User ID and Track ID, provided a strong foundation for the FM model, as evidenced by the MRR scores. The inclusion of Valence, a measure of musical positiveness, led to a notable increase in MRR in both POP RND and POP USER settings, underscoring its predictive power for user preference. The addition of Tempo, while resulting in a lower MRR compared to the baseline+Valence model, still demonstrated the value of rhythmic characteristics in influencing user preferences.

Methods involving contextual data, such as hashtags (Method 4) and sentiment analysis (Method 5), showed different effects in the two settings. While the inclusion of hashtags led to a considerable MRR increase in the POP USER setting, it was less impactful in the POP RND setting. This discrepancy may indicate that hashtags provide valuable context clues when users are already familiar with the tracks, making them more influential in discerning the user's current mood or situation.

Sentiment analysis exhibited the highest MRR in the POP USER setting, suggesting that the emotional content inferred from user tweets plays a significant role when users choose tracks within familiar repertoires. This also hints at the potential of sentiment analysis as a proxy for complex emotional states that influence listening behavior, particularly in personalized settings.

Our evaluation through Mean Reciprocal Rank (MRR) unveils a nuanced understanding of the factors influencing music recommendation systems using Factorization Machines (FMs). This discussion interprets the results, considering the implications of content and contextual features across different user settings.

7.1.1 Influence of Features Across Settings. The baseline FM method, employing only user and track identifiers, achieved a relatively high MRR in the POP RND setting but was less effective in the POP USER scenario. This suggests that while basic user and track ID data can predict preferences across random selections, it may not suffice when users are familiar with the tracks, emphasizing the need for more personalized features in such settings.

7.1.2 Content Features and Cold Start Problems. The integration of track-specific content features, such as Valence and Tempo, resulted in enhanced MRR for the POP RND setting. This improvement indicates the value of these features in scenarios resembling cold start problems, where little is known about user preferences, and recommendations need to leverage the inherent characteristics of the music tracks themselves.

7.1.3 Impact of Contextual Data in Personalized Settings. The inclusion of hashtags and sentiment analysis derived from them substantially improved MRR results in the POP USER setting, illustrating their effectiveness in personalized recommendation scenarios. These features appear to offer significant clues about user preferences that are more discernible when the users are engaged with known tracks. In contrast, their contribution is less pronounced in the POP RND setting, suggesting that the role of such contextual data is contingent upon the user's prior familiarity with the music.

7.1.4 Content Features in User-Based Settings. Notably, while content features such as Valence and Tempo contribute positively to the results in the POP RND setting, their impact is notably reduced in the personalized POP USER setting. This could indicate that when users are choosing from a selection of familiar tracks, content features alone are insufficient to predict preferences accurately, and the system benefits from the inclusion of contextual information that captures the user's current state or environment.

7.2 Implications for Context-Aware Recommendations

The results suggest that context-aware features, especially sentiment-related ones, are more effective in personalized environments where users are already acquainted with the tracks. This aligns with the hypothesis that recommendations can be significantly improved by understanding not just historical preferences but also the current contextual state of the user.

7.3 Challenges and Considerations

However, the lower MRR scores observed for some feature integrations highlight the challenges in context-aware recommendation systems. For instance, the relative drop in MRR when incorporating features such as Tempo alone suggests that not all content features equally contribute to the recommendation quality, and their effectiveness might depend on the combination with other contextual clues.

The study also noted the limitation of FMs in terms of capturing self-interaction terms, which may result in information loss. Although the methodology attempted to mitigate this by including a wide range of feature terms, the impact of this limitation on the recommendation quality warrants further investigation.

7.4 Future Directions

These observations lead to several future research directions. For instance, exploring a broader set of contextual features, such as specific activity indicators or environmental variables like weather, could further enhance recommendation relevance. Additionally, adjusting the FM algorithm to capture more complex interactions or integrating it with deep learning models may offer ways to improve recommendation quality, even under high data sparsity.

In summary, the discussion reveals that while context-aware features enhance the FM-based recommendation system's performance, the selection and integration of these features are critical. The findings from this study provide a solid foundation for future work aimed at refining the personalization capabilities of music recommendation systems.

8 CONCLUSION

This study presented a comprehensive evaluation of a Context-Aware Music Recommendation System utilizing Factorization Machines (FMs). Through rigorous experiments using the #nowplaying-RS dataset, we demonstrated the significant role of contextual and content features in enhancing the accuracy of music recommendations.

Our findings indicate that while the base FM model achieves considerable success in general settings such as POP RND, it falls short in more personalized contexts, as seen in the POP USER setting. The addition of content features like Valence and Tempo improves performance in addressing cold start issues, but their influence is less dominant in personalized scenarios where users have established preferences.

Contextual features derived from hashtags and sentiment analysis emerge as powerful enhancers of personalization, significantly improving MRR scores in settings where users are familiar with the tracks. This underscores the importance of mood and context-specific information in developing sophisticated recommendation systems.

In conclusion, the #nowplaying-RS dataset serves as a valuable benchmark for context-aware recommendation systems, facilitating the exploration of innovative solutions that blend technological sophistication with practical applicability. The insights gained lay the groundwork for future enhancements, including the integration of more diverse contextual data and the development of advanced algorithms to capture complex user preferences. As music listening becomes increasingly integrated into daily life, the pursuit of a nuanced, user-centered approach to music recommendation continues to be a compelling and essential endeavor.

REFERENCES

[1] Gediminas Adomavicius and Alexander Tuzhilin. 2010. Context-aware recommender systems. In *Recommender systems handbook*. Springer, 217–253.

[2] Marissa Baxter, Lisa Ha, Kirill Perfiliev, and Natalie Sayre. 2021. Context-Based Music Recommendation Algorithm Evaluation. *arXiv preprint arXiv:2112.10612* (2021).

[3] Jia-Wei Chang, Ching-Yi Chiou, Jia-Yi Liao, Ying-Kai Hung, Chien-Che Huang, Kuan-Cheng Lin, and Ying-Hung Pu. 2021. Music recommender using deep embedding-based features and behavior-based reinforcement learning. *Multimedia Tools and Applications* (2021), 1–28.

[4] Ahmet Elbir and Nizamettin Aydin. 2020. Music genre classification and music recommendation by using deep learning. *Electronics Letters* 56, 12 (2020), 627–629.

[5] Martin Pichl and Eva Zangerle. 2021. User models for multi-context-aware music recommendation. *Multimedia Tools and Applications* 80, 15 (2021), 22509–22531.

[6] Martin Pichl, Eva Zangerle, and Günther Specht. 2015. Towards a context-aware music recommendation approach: What is hidden in the playlist name?. In *2015 IEEE international conference on data mining workshop (ICDMW)*. IEEE, 1360–1365.

[7] Asmita Poddar, Eva Zangerle, and Yi-Hsuan Yang. 2018. nowplaying-RS: a new benchmark dataset for building context-aware music recommender systems. In *Proceedings of the 15th Sound & Music Computing Conference*. 21–26.

[8] Steffen Rendle. 2010. Factorization machines. In *2010 IEEE International conference on data mining*. IEEE, 995–1000.

[9] Steffen Rendle. 2012. Factorization Machines with libFM. *ACM Trans. Intell. Syst. Technol.* 3, 3, Article 57 (may 2012), 22 pages. https://doi.org/10.1145/2168752.2168771

[10] Hamed Zamani, Markus Schedl, Paul Lamere, and Ching-Wei Chen. 2019. An analysis of approaches taken in the acm recsys challenge 2018 for automatic music playlist continuation. *ACM Transactions on Intelligent Systems and Technology (TIST)* 10, 5 (2019), 1–21.

[11] Yong Zheng and David Wang. 2022. Multi-criteria ranking: Next generation of multi-criteria recommendation framework. *IEEE Access* 10 (2022), 90715–90725.